

# ARUSHI JAIN

Mila, McGill University, Montreal, Canada

[arushi.jain@mail.mcgill.ca](mailto:arushi.jain@mail.mcgill.ca) ♦ <https://arushijain94.github.io/>

## RESEARCH INTERESTS

---

Reinforcement Learning (RL), Safe RL, Off-Policy RL, Constrained MDPs.

## EDUCATION

---

### McGill University, Montreal, Canada

Ph.D., Computer Science	GPA: 4.00/4.00	<i>Sept 2019 - Present</i>
MSc, Computer Science	GPA: 4.00/4.00	<i>Sept 2017 - Sept 2019</i>

### Mila

Supervisor: [Doina Precup](#), [Pierre-Luc Bacon](#)

### Indraprastha Institute of Information Technology Delhi (IIIT-D), India

Bachelor of Technology, Computer Science and Engineering	GPA: 9.42/10.00	<i>Aug 2012 - May 2016</i>
--	-----------------	----------------------------

Image Analysis and Biometrics (IAB) lab

Supervisor: [Mayank Vatsa](#), [Richa Singh](#)

## AWARDS

---

- **Bourse du Centre de Recherche Informatique de Montréal (CRIM) pour Études Supérieures (2018-2019)** Graduate Scholarship provides financial support to only one master and a doctoral student.

## PUBLICATIONS

---

1. **Towards Painless Policy Optimization for Constrained MDPs**[\[Paper\]](#)  
**Arushi Jain**, Sharan Vaswani, Reza Babanezhad, Csaba Szepesvari, Doina Precup  
Accepted in Conference on Uncertainty in Artificial Intelligence (UAI), 2022.  
Accepted at RLDM 2022 [\[Paper\]](#)
2. **Variance Penalized On-Policy and Off-Policy Actor-Critic**[\[Paper\]](#)[\[Poster\]](#)[\[Slide\]](#)  
**Arushi Jain**, Gandharv Patil, Ayush Jain, Khimya Khetarpal, Doina Precup  
Accepted to the Thirty-Fifth AAAI Conference on Artificial Intelligence (AAAI), 2021.
3. **Safe Option-Critic: Learning Safety in the Option-Critic Architecture**[\[Paper\]](#)[\[Poster\]](#)[\[Slide\]](#)  
**Arushi Jain**, Khimya Khetarpal, and Doina Precup  
Accepted in Adaptive Learning Agents (ALA), ICML Workshop, 2018.  
Accepted in **The Knowledge Engineering Review (KER) Journal** 2021. [Impact Factor 1.07]
4. **Safe Actor-Critic**[\[Paper\]](#)[\[Poster\]](#)[\[Slide\]](#)  
**Arushi Jain**, Ayush Jain, and Doina Precup  
Accepted in Safety, Risk and Uncertainty in RL, Uncertainty in AI (UAI) Workshop, 2018.  
Accepted in Women in ML (WiML), (NeurIPS) Workshop, 2018.
5. **Learning Options using Constrained Return Variance**[\[Paper\]](#)[\[Poster\]](#)  
**Arushi Jain** and Doina Precup  
Accepted in Safety and Robustness in Decision Making, NeurIPS Workshop 2019.

6. **Safe Policy Learning with Constrained Return Variance**[Paper][Poster]

**Arushi Jain**

Accepted in Graduate Student AI Symposium held at **Canadian AI Conference** 2019. The proceeding is also published in the **LNAI** Series by Springer Verlag.

7. **Safe Hierarchical Policy Optimization using Constrained Return Variance in Options**[Paper]

**Arushi Jain**, Doina Precup

Accepted at **RLDM** 2019.

8. **Safety using constraint variance in policy-gradient methods**[Thesis]

**Arushi Jain**

Master Thesis, McGill University, Montreal, March 2020.

## RESEARCH & WORK EXPERIENCE

---

**Mila, RL Lab, McGill University**

September 2017 - Present

*Graduate Student* advised by *Prof. Doina Precup, Pierre-Luc Bacon*

- Real-world tasks with different competing objectives or constraints, it is not always feasible to evaluate the performance of a given policy/behavior. For example, one wants to assess the performance of a given policy in recommendation systems. However, allowing a potentially bad policy that does not abide system's constraints to interact with customers can lead to unpredictable losses. To overcome this problem, we propose to predict by learning from a different **off-policy** that follows constraints and also leads to smaller variability in the performance than the standard on-policy approach.
- We designed a novel and generic **on-policy and off-policy safe actor-critic** framework to learn safe policies such that objective function trades off between minimizing the effects of model uncertainty and maximizing the mean return.
- We worked on introducing novel safety solutions in Options (**hierarchical RL**) by deriving a policy-gradient algorithm where variance in the temporal difference (TD) error is regularized.

**Mila, Amii**

August 2021 - March 2022

*Mentored by Sharan Vaswani, Reza Babanezhad, Csaba Szepesvri, Doina Precup*

We study planning in infinite horizon Constrained MDPs, where we use parameter-free algorithm from online linear optimization to learn both the policy and follow the constraints. Unlike gradient descent-ascent methods, it is robust to hyperparameter tuning and environment misspecification.

**SPORTLOGiQ, Montreal, Canada**

June 2019 - Sept 2019

*RL Research Intern* mentored by *Norm Ferns*

Worked on formally comparing and contrasting the agents in a given Markov Decision Process (MDP) by coming up with behavioral pseudo-metrics following the work on *lax-bisimulation*.

**Borealis AI, Edmonton, Canada**

May 2018 - Aug 2018

*AI Research Intern* mentored by *Nidhi Hedge*

Worked on **safe recommendations** in RL setting which would help provide recommendations with certain guarantees. This would provide life-long learning as well as long-term goal maximization which would be particularly useful in financial applications.

**Microsoft Research (MSR), India**

June 2016 - July 2017

*Research Fellow* advised by *Sundararajan Sellamanickam, Arun Iyer*

- Developed a **service monitoring and diagnostic tool** for unsupervised hierarchical monitoring of services which replaced a heuristic-based system failing to jointly model multivariate time-series.

- Worked on **email solutions using DNN** to build intelligent approaches to solve problems like multi-label classification, learning a useful representation of email content, etc.

## TECHNICAL SKILLS & COURSEWORK

---

**Programming Languages:** Python, Tensorflow, Pytorch, C

**Coursework:** Reinforcement Learning, Probabilistic Analysis of Algorithms, Applied ML, Theoretical ML, Probabilistic Graphical Modeling, Matrix Computation, Mathematical Foundation of ML, Reinforcement Learning and Optimal Control.